

Medicine Outcome Prediction System

LIME model explainability | Held-out test-set interpretation

Purpose

LIME explains how each feature pushes the model output toward a Yes treatment outcome on a small held-out sample of the test set.

Interpretation

Contribution scores are local explanations for individual patients and should be interpreted as model influence, not causation. Feature importance can reflect correlations, missingness, and class imbalance in the data.

Analysis settings


Models: 4

Background rows: 25


Explained rows per model: 20

Target explained: positive-class probability (treatment_outcome = Yes)


Decision Tree: LIME feature importance on held-out test data



Random Forest: LIME feature importance on held-out test data



XGBoost: LIME feature importance on held-out test data



Gradient Boosting: LIME feature importance on held-out test data

